









Evaluating a Learning Algorithm

Bias vs. Variance

	Diagnosing Bias vs. Variance	7 min
	Diagnosing Bias vs. Variance	3 min
	Regularization and Bias/Variance	11 min
	Regularization and Bias/Variance	3 min
	Learning Curves	11 min
	Learning Curves	3 min
	Deciding What to Do Next Revisited	6 min
	Deciding What to do Next Revisited	3 min

Review

Building a Spam Classifier

Handling Skewed Data

Using Large Data Sets

Review

Deciding What to Do Next Revisited

Our decision process can be broken down as follows:

- **Getting more training examples:** Fixes high variance
- **Trying smaller sets of features:** Fixes high variance
- **Adding features:** Fixes high bias
- **Adding polynomial features:** Fixes high bias
- **Decreasing λ :** Fixes high bias
- **Increasing λ :** Fixes high variance.

Diagnosing Neural Networks

- A neural network with fewer parameters is **prone to underfitting**. It is also **computationally cheaper**.
- A large neural network with more parameters is **prone to overfitting**. It is also **computationally expensive**. In this case you can use regularization (increase λ) to address the overfitting.

Using a single hidden layer is a good starting default. You can train your neural network on a number of hidden layers using your cross validation set. You can then select the one that performs best.

Model Complexity Effects:

- Lower-order polynomials (low model complexity) have high bias and low variance. In this case, the model fits poorly consistently.
- Higher-order polynomials (high model complexity) fit the training data extremely well and the test data extremely poorly. These have low bias on the training data, but very high variance.
- In reality, we would want to choose a model somewhere in between, that can generalize well but also fits the data reasonably well.

Marcar como completo

